

CS516: Parallelization of Programs

Introduction

Vishwesh Jatala

Assistant Professor

Department of CSE

Indian Institute of Technology Bhilai

vishwesh@iitbhilai.ac.in



2023-24 W

What is the output?


```
1#include<stdio.h>
2
3void main()
4{
5
6    printf("-----\n\n\t");
7    printf("%c%c%c%c%c%c%c%c%c%c%c%c%c%c%c%c%c%c%c",
8          72,65,80,80,89,32,78,69,87,32,89,69,65,82,32,50,48,50,52,33);
9    printf("\n\n-----\n");
10}
```

Outline of Today's Lecture


- Why?
- What?
- How?

Moore's Law



- *The number of transistors on a IC doubles about every two years*



Moore's Law Effect





Moore's Law Effect



Single Core Performance




Power and Heat

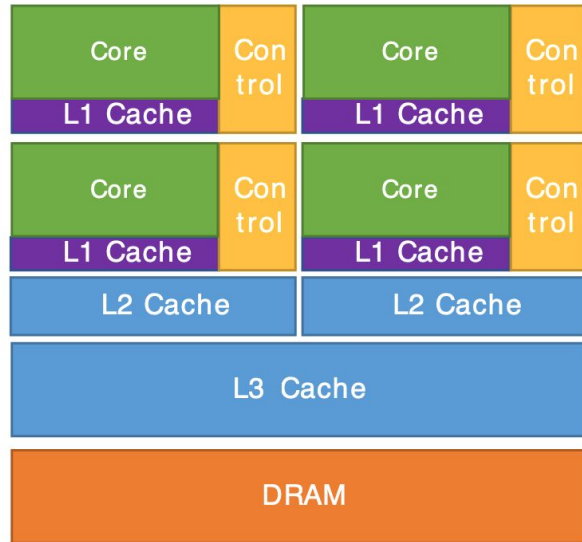


Multicore Processors


Parallel Architectures are Everywhere!



Parallel Hardware



Multicores



Distributed CPUs

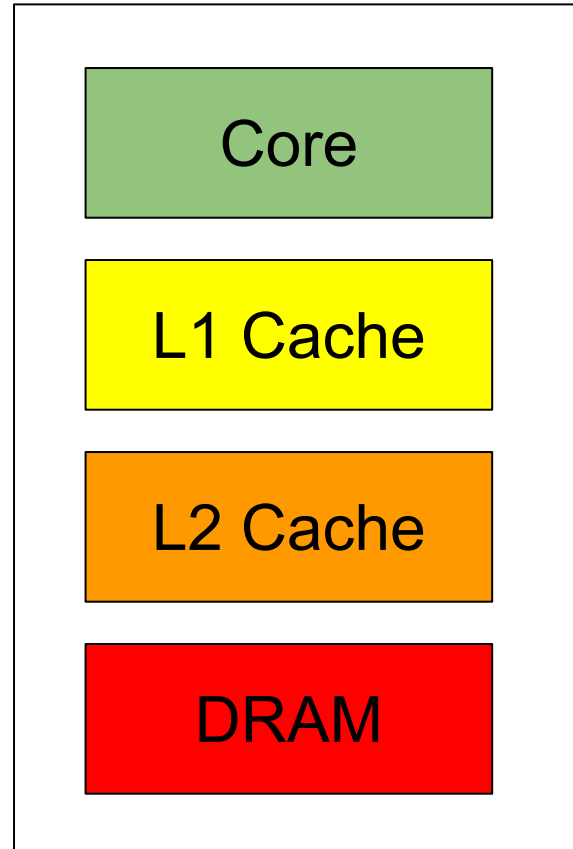


GPUs

Hardware and Software

```
1 package com.beginnersbook;
2 public class JavaExample {
3     public static void main(String []args) {
4         String str[] = { "Ajeet", "Steve", "Rick", "I"
5         String temp;
6         System.out.println("Strings in sorted order:");
7         for (int j = 0; j < str.length; j++) {
8             for (int i = j + 1; i < str.length; i++)
9                 // comparing adjacent strings
10                if (str[i].compareTo(str[j]) < 0) {
11                    temp = str[j];
12                    str[j] = str[i];
13                    str[i] = temp;
14                }
15            }
16            System.out.println(str[j]);
17        }
18    }
19 }
```

Sequential



Single-core CPU

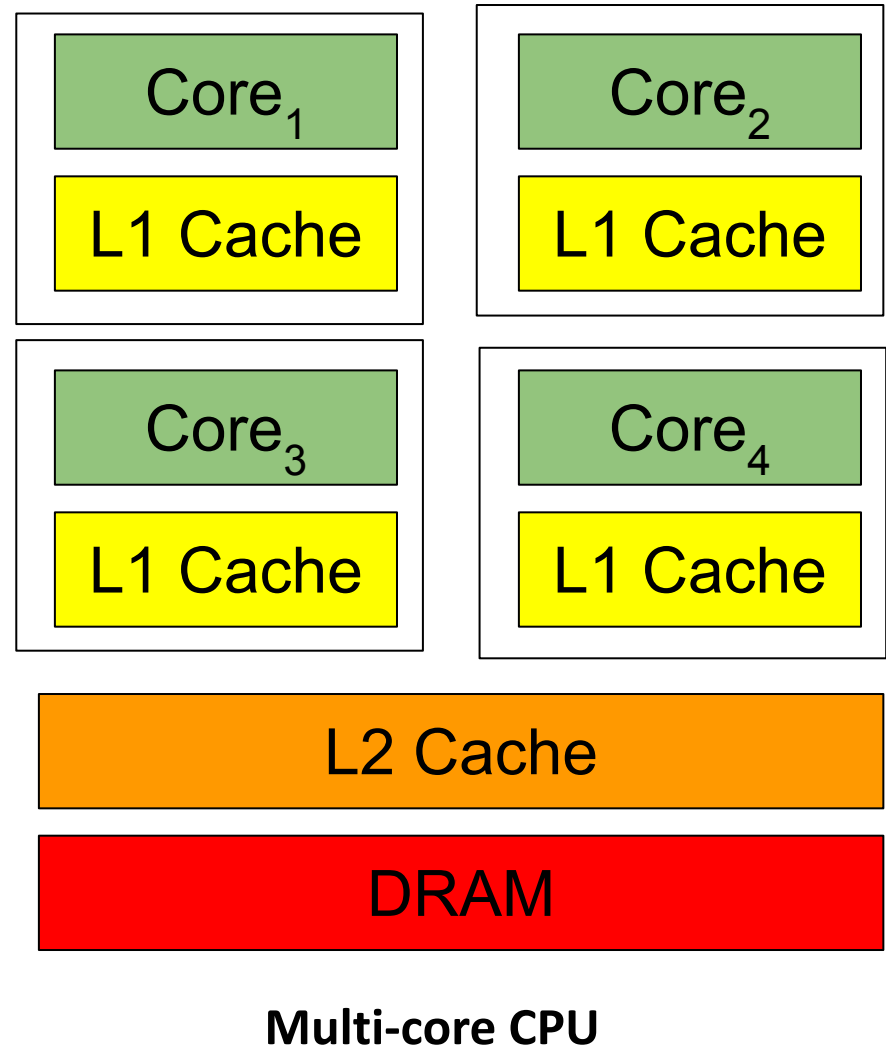


Output

Hardware and Software


```
1 package com.beginnersbook;
2 public class JavaExample {
3     public static void main(String []args) {
4         String str[] = { "Ajeet", "Steve", "Rick", "I"
5         String temp;
6         System.out.println("Strings in sorted order:");
7         for (int j = 0; j < str.length; j++) {
8             for (int i = j + 1; i < str.length; i++)
9                 // comparing adjacent strings
10                if (str[i].compareTo(str[j]) < 0) {
11                    temp = str[j];
12                    str[j] = str[i];
13                    str[i] = temp;
14                }
15            }
16            System.out.println(str[j]);
17        }
18    }
19 }
```

Same sequential?



Professor P

15 questions
300 Answer sheets



Professor P's Teaching Assistants



Benefits of Parallel Programming



Fast: Less
execution time




Save Money



Solves Larger Problem


Parallel Programming Applications




Deep Learning and
Machine Learning



Medical Imaging



CFD



Climate Modeling

Parallel Programming Applications

OpenAI used a super computer
with more than 285,000 CPU cores,
10,000 GPUs and
400 gigabits per second of network
connectivity for each GPU server.

Challenges!

Example-1

Sequential Execution

```
for (int i=0; i<5; i++)  
    A[i]= i
```

Core-0



A[0]=0

A[1]=1

A[2]=2

A[3]=3

A[4]=4

Example-1

```
for (int i=0; i<5; i++)  
    A[i]= i
```


Parallel Execution

Core-0




A[0]=0

Core-1




A[1]=1

Core-2



A[2]=2

Core-3



A[3]=3

Core-4




A[4]=4

Example-2

Sequential Execution

```
int count = 0;  
for (int i=0; i<5; i++)  
    A[i]= count++;
```

Core-0



A[0]=0

A[1]=1

A[2]=2

A[3]=3

A[4]=4

Example-2

```
int count = 0;  
for (int i=0; i<5; i++)  
    A[i]= count++;
```


Parallel Execution

Core-0




A[0]=1

Core-1




A[1]=2

Core-2




A[2]=0

Core-3



A[3]=4

Core-4



A[4]=3

Challenges:



Detecting Parallelism is Hard!

Example-3: Sequential Version

```
int sum = 0;
for (i = 0; i < n; i++) {
    x = f(i);
    sum = sum+x;
}
```

A Sequential Program for Sum

Core-0


1	4	3	9	2	8	5	1	1	6	2	7	2	5	0	4	1	8	6	5	1	2	3	9
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

sum = 95

Example-3: Parallel Version-1

```
my_sum = 0;  
my_first i = . . . ;  
my_last i = . . . ;  
for (my_i = my_first i; my_i < p_last_i; my_i++) {  
    my_x = f(i);  
    my_sum += my_x;  
}
```

Core-0 Core-1 Core-2 Core-3 Core-4 Core-5 Core-6 Core-7




my_sum 8 19 7 15 7 13 12 14

Example-3: Parallel Version-1

```
if (I'm the master core) {  
    sum = my sum;  
    for (each core other than myself) {  
        receive value from core;  
        sum += value;  
    }  
}  
else  
    { send my_sum to the master; }
```


Core-0 Core-1 Core-2 Core-3 Core-4 Core-5 Core-6 Core-7



my_sum 8 19 7 15 7 13 12 14

sum 95 (8+19+7+15+7+13+12+14)

Example-3: Parallel Version-2



Parallel Version-1 or Version-2?

- Both have same number of operations
- Version-1 sum is sequential
- Version-2 exposes parallelism

Challenges:

Detecting Parallelism is Hard!

Communication

Synchronization

Example-4:

```
int sum = 0;
for (i = 0; i < n; i++) {
    x = f(i);
    sum = sum+x;
}
```

Challenges:

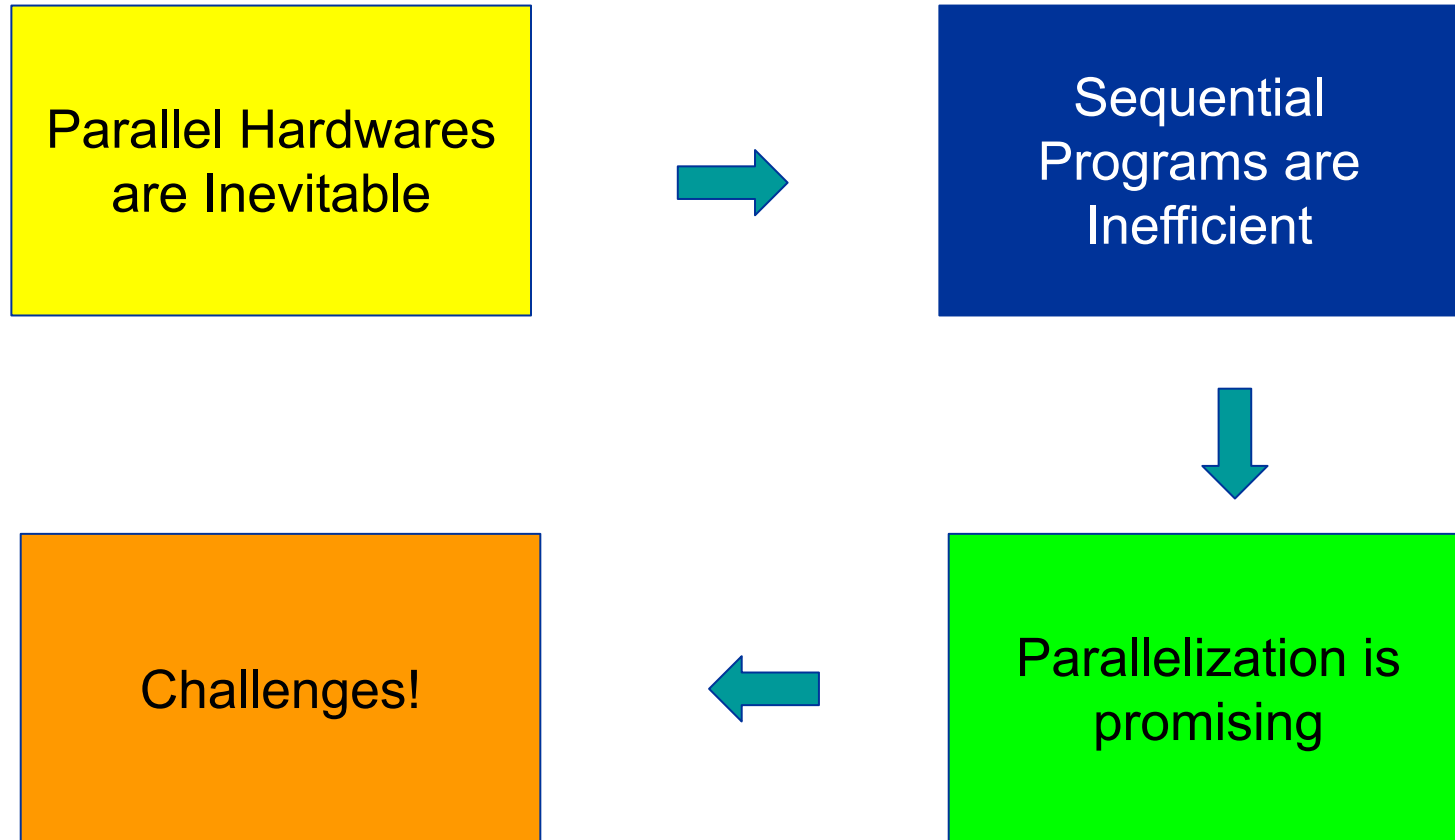
Detecting Parallelism is Hard!

Communication

Synchronization

Load Balancing

What did we learn so far?





Let's discuss details from next class!

Course Outline (Part-1)

- Introduction (this lecture)
- Overview of Parallel Architectures and Programming models
- Amdahl's law and Performance
- Parallel programming
 - GPUs and CUDA programming
 - Optimizations
- Case studies
- Extracting Parallelism from Sequential Programs Automatically

Course Logistics

- Lecture Hours:
 - Mon, Tue, Thursday 10:30 am - 11:25 am
- Course Website: Canvas platform
 - Lecture notes
 - Assignments
 - Project
 - Discussions
 - Marks

Course Logistics: Evaluation Scheme

- Evaluation scheme (can be changed slightly):
 - ❑ Exams: ~40%
 - ❑ Project: ~35%
 - ❑ Attendance: ~10%
 - ❑ Assignments (Paper presentation?): 15%

- Attendance
 - ❑ 0% - 50%: 0 Marks
 - ❑ >50%: Marks will be awarded out of 10 accordingly.
 - ❑ Example:
 - Total sessions: 16
 - #sessions attended = 7 (<50%), marks = 0
 - #sessions attended = 10 (62.5%), marks = 2.5 ($2 \times 10/8$)

Course Logistics: References

- Lecture notes will be available on Canvas
- Reference material will be provided on Canvas
- Text book for extracting parallelism:
 - Randy Allen, Ken Kennedy, Optimizing Compilers for Modern Architectures: A Dependence-based Approach, Morgan Kaufmann, 2001

Course Logistics: Tools

- Platforms:
 - ❑ Prefer Google Colab for GPUs and CUDA Programming.
 - ❑ A demo session

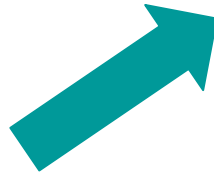
Course Logistics

- Evaluation Policy:
 - Acknowledge all the sources
 - **Do not cheat**

Outcome of the Course?

- State-of-the-art techniques in parallel computing
- Develop parallel programming skills
- Transferable skills -
 - Parallel programming is used in multiple disciplines
 - Industries
 - Education and research
- Handle projects
- High-performance computing

Course Logistics



Office: B-010



vishwesh@iitbhilai.ac.in



CANVAS

References

- <https://www.cse.iitd.ac.in/~soham/COL380/page.html>
- <https://s3.wp.wsu.edu/uploads/sites/1122/2017/05/6-9-2017-slides-vFinal.pptx>
- Miscellaneous resources on internet



Thank you!