

CS516: Parallelization of Programs

CUDA Thread Organization

Vishwesh Jatala

Assistant Professor

Department of EECS

Indian Institute of Technology Bhilai

vishwesh@iitbhilai.ac.in

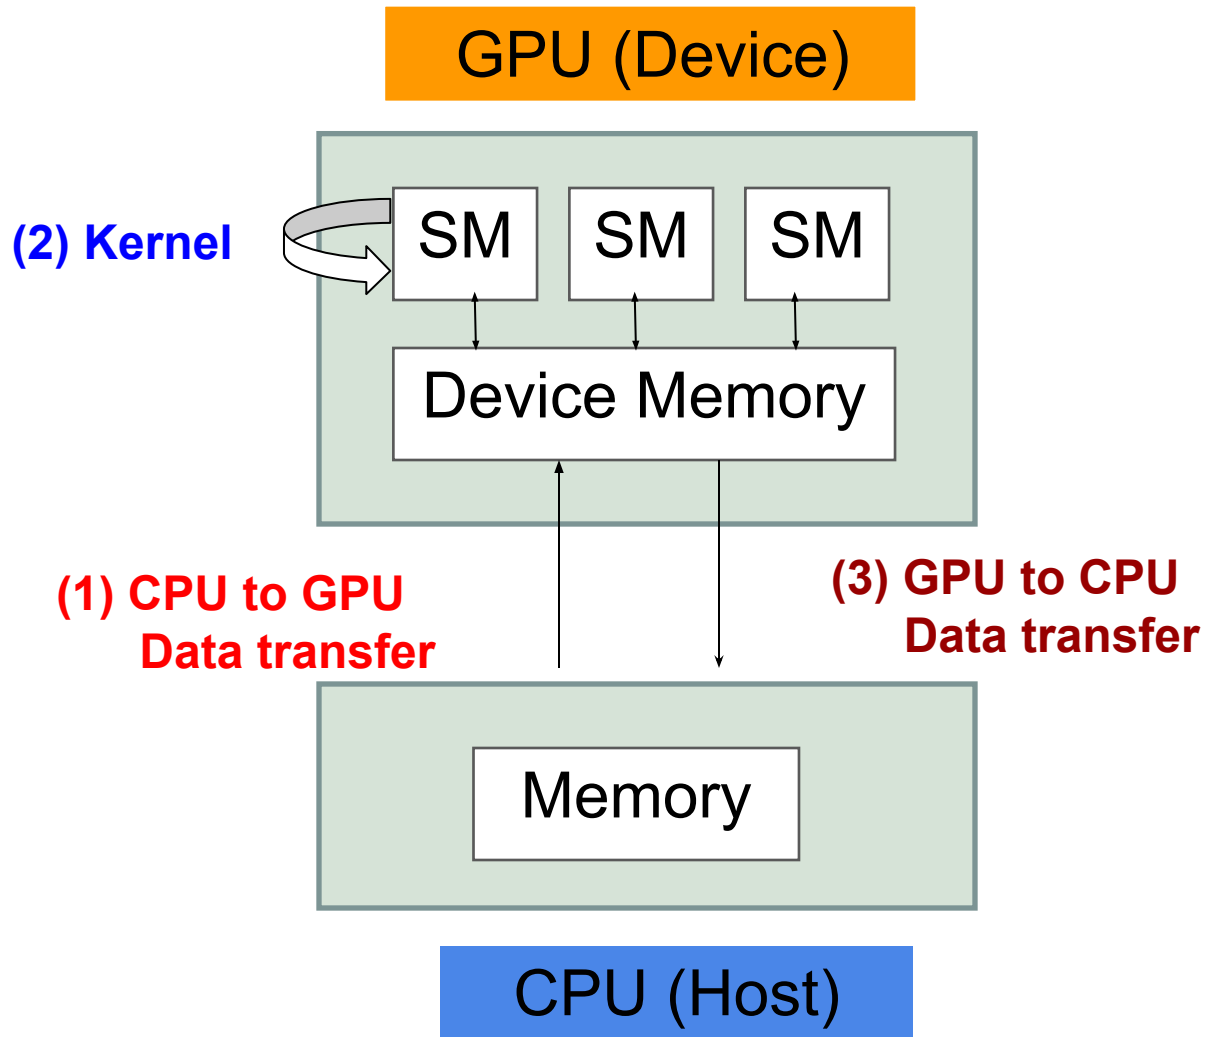


2022-23M

Outline

- Continue with CUDA Programming
 - Thread organization
 - Examples

CUDA Programming Flow



VectorAdd in CUDA

- For given two vectors A and B both having size N (where $N \leq 1024$), write a CUDA program to compute $C=A+B$

VectorAdd in CUDA

```
#define N 512
int main(void) {
    int *a, *b, *c; // host copies of a, b, c
    int *dev a, *dev b, *dev c; //device copies of a, b, c
    int size = N * sizeof(int);

    ...
    // Alloc space for host copies of a, b, c and
    // setup input values
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);

    // Alloc space for device copies of a, b, c
    cudaMalloc((void **)&dev a, size);
    cudaMalloc((void **)&dev_b, size);
    cudaMalloc((void **)&dev_c, size);
```

VectorAdd in CUDA

```
// Copy inputs to device
cudaMemcpy(dev_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(dev_b, b, size, cudaMemcpyHostToDevice);

// Launch add() kernel on GPU with N threads
add<<<1,N>>>(dev_a, dev_b, dev_c);

// Copy result back to host
cudaMemcpy(c, dev_c, size, cudaMemcpyDeviceToHost);
// Cleanup
free(a); free(b); free(c);
cudaFree(dev_a); cudaFree(dev_b); cudaFree(dev_c);
return 0;
}
```

VectorAdd in CUDA

```
__global__ void add(int *a, int *b, int *c) {  
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];  
}
```

Practise Problem-1


- For given two matrices M and N both having size $k \times k$ (where $k \leq 1024$), write a CUDA program to perform $M+N$
 - *Hint: Allocate M and N single dimension array having $k \times k$ elements.*

Thread Configuration

```
add<<<ThreadConfig>>> (dev_a, dev_b, dev_c);
```



```
add<<<ThreadBlocks, Threads>>> (dev_a, dev_b, dev_c);
```

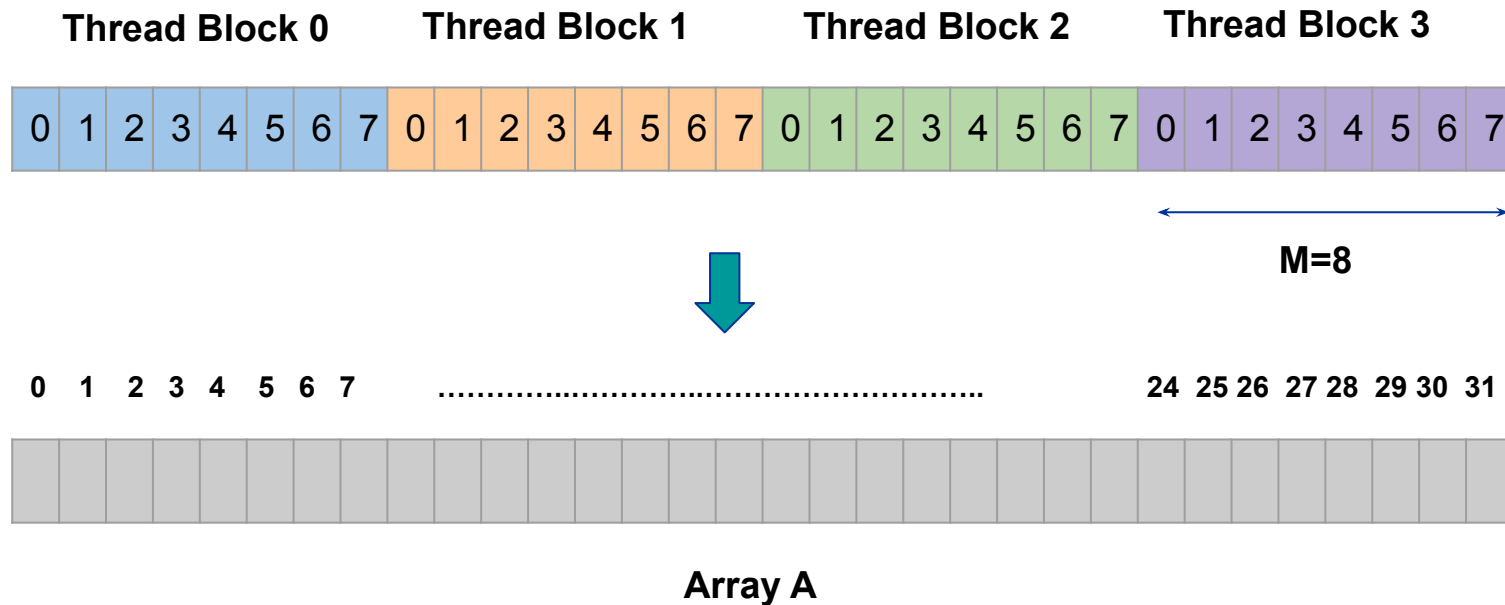


`blockIdx.x` -> For identifying the block

`threadIdx.x` -> For identifying the thread within a thread block

`blockDim.x` -> Size of thread block


Indexing Arrays with Threads and Thread Blocks



What is the array index accessed by thread having `threadIdx.x` from the `blockIdx.x`?

```
int index = threadIdx.x + blockIdx.x * M;
```

Indexing Arrays with Threads and Thread Blocks



Which **threadIdx.x** and **BlockIdx.x** will operate on index **21**?

$$\begin{aligned} \text{index} &= \text{threadIdx.x} + \text{blockIdx.x} * M; \\ 21 &= 5 + 2 * 8 \end{aligned}$$

VectorAdd in CUDA with Thread and Blocks

```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[index] = a[index] + b[index];  
}
```

VectorAdd in CUDA

```
#define N 512
int main(void) {
    int *a, *b, *c; // host copies of a, b, c
    int *dev a, *dev b, *dev c; //device copies of a, b, c
    int size = N * sizeof(int);

    ...
    // Alloc space for host copies of a, b, c and
    // setup input values
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);

    // Alloc space for device copies of a, b, c
    cudaMalloc((void **)&dev a, size);
    cudaMalloc((void **)&dev_b, size);
    cudaMalloc((void **)&dev_c, size);
```

VectorAdd in CUDA

```
// Copy inputs to device
cudaMemcpy(dev_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(dev_b, b, size, cudaMemcpyHostToDevice);

// Launch add() kernel on GPU with N threads
add<<<<N/THREADS_PER_BLOCK,THREADS_PER_BLOCK>>>>
    (dev_a, dev_b, dev_c);

// Copy result back to host
cudaMemcpy(c, dev_c, size, cudaMemcpyDeviceToHost);
// Cleanup
free(a); free(b); free(c);
cudaFree(dev_a); cudaFree(dev_b); cudaFree(dev_c);
return 0;
}
```

Threadblock configuration

```
add<<<ThreadBlocks, Threads>>> (dev_a, dev_b, dev_c);
```

- Thread block configuration
 - ❑ User choice
 - ❑ Depends on problem size
- Problem size = 32768 ($1024 * 32$)
 - ❑ Threadblocks = 32, No of threads/thread block = 1024
 - ❑ Threadblocks = 128, No of threads/thread block = 256

Practise Problem-2


- For given two matrices M and N both having size $k \times k$, write a CUDA program to perform $M+N$ using threads and thread blocks
 - Assume threads in a thread block as `THREADS_PER_BLOCK`

Thread Configuration

```
add<<<ThreadConfig>>> (dev_a, dev_b, dev_c);
```



```
add<<<ThreadBlocks, Threads>>> (dev_a, dev_b, dev_c);
```



Exercise: 1D Thread Organization

- For a given matrix A of size $N \times M$, write a CUDA program to initialize the matrix elements as below

0	1	2	3	4	5
6	7	8	9	10	11
12	13	14	15	16	17
18	19	20	21	22	23
24	25	26	27	28	29

- **Assumptions:**
 - ❑ Matrix is stored in the single dimensional array.
 - ❑ No. of thread blocks = N
 - ❑ No. of threads in each block = M

1D

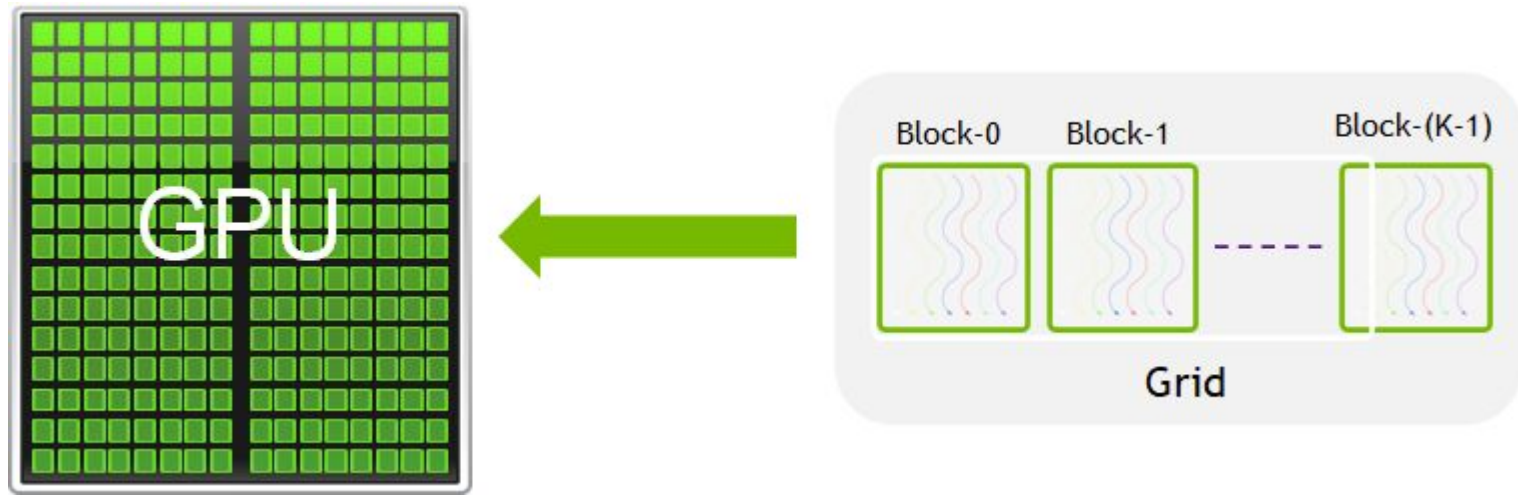
```
#include <stdio.h>
#include <cuda.h>
__global__ void dkernel(unsigned *matrix) {
    unsigned id = blockIdx.x * blockDim.x + threadIdx.x;
    matrix[id] = id;
}
#define N    5
#define M    6
int main() {
    unsigned *matrix, *hmatrix;

    cudaMalloc(&matrix, N * M * sizeof(unsigned));
    hmatrix = (unsigned *)malloc(N * M * sizeof(unsigned));

    dkernel<<<N, M>>>(matrix);
    cudaMemcpy(hmatrix, matrix, N * M * sizeof(unsigned), cudaMemcpyDeviceToHost);

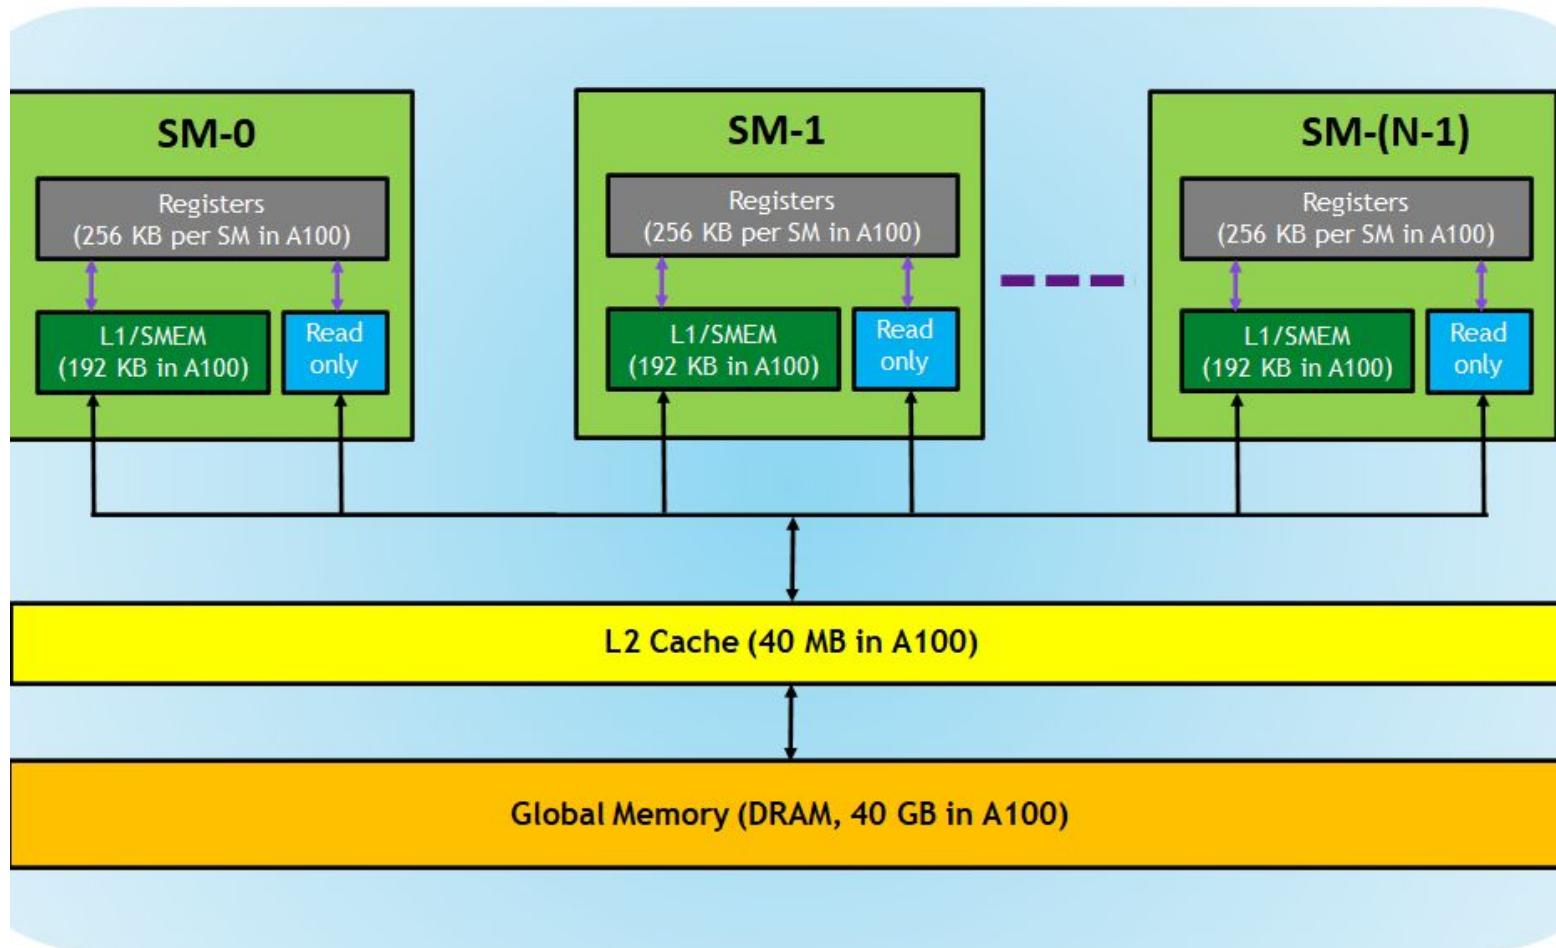
    for (unsigned ii = 0; ii < N; ++ii) {
        for (unsigned jj = 0; jj < M; ++jj) {
            printf("%2d ", hmatrix[ii * M + jj]);
        }
        printf("\n");
    }
    return 0;
}
```

Why Thread Blocks?



But why not 1 thread block with all the threads in it?

GPU Architecture




Few Constraints

- Thread block size has limit
 - ❑ Each thread executes the same program
 - ❑ Each thread requires some registers
 - ❑ Number of registers in each SM is finite



```
__global__ void add(int *a, int *b, int *c) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[index] = a[index] + b[index];  
}
```

Scalability


GPU-0




GPU-1



SM₀



SM₁



SM₀

SM₁

SM₂

SM₃

Few Constraints

- Thread block size has limit
- Max number of threads blocks reside per SM
- Max threads reside per SM

Compute Capabilities

GPU	Kepler GK180	Maxwell GM200	Pascal GP100	Volta GV100
Compute Capability	3.5	5.2	6.0	7.0
Threads / Warp	32	32	32	32
Max Warps / SM	64	64	64	64
Max Threads / SM	2048	2048	2048	2048
Max Thread Blocks / SM	16	32	32	32
Max 32-bit Registers / SM	65536	65536	65536	65536
Max Registers / Block	65536	32768	65536	65536
Max Registers / Thread	255	255	255	255 ¹
Max Thread Block Size	1024	1024	1024	1024
FP32 Cores / SM	192	128	64	64
Ratio of SM Registers to FP32 Cores	341	512	1024	1024
Shared Memory Size / SM	16 KB/32 KB/ 48 KB	96 KB	64 KB	Configurable up to 96 KB

Source: <https://images.nvidia.com/content/volta-architecture/pdf/volta-architecture-whitepaper.pdf>

Summary

- CUDA Programming
 - Thread organizations:
 - Examples
- Next Lecture
 - Thread organization (2D & 3D)
 - GPU Instruction Execution

References

- CS6023 GPU Programming
 - <https://www.cse.iitm.ac.in/~rupesh/teaching/gpu/jan20/>
- Miscellaneous resources from internet