

CS516: Parallelization of Programs

GPU Memories

Vishwesh Jatala

Assistant Professor

Department of CSE

Indian Institute of Technology Bhilai

vishwesh@iitbhilai.ac.in




2022-23 W

References

- Miscellaneous resources from internet
- CS6023 GPU Programming
 - <https://www.cse.iitm.ac.in/~rupesh/teaching/gpu/jan20/>


Shared Memory

- Shared memory is faster but limited in size
- To improve concurrency, shared memory is divided into banks



Shared Memory Banks

- A memory address is translated to a bank
- Addresses that fall to different banks can be accessed in parallel

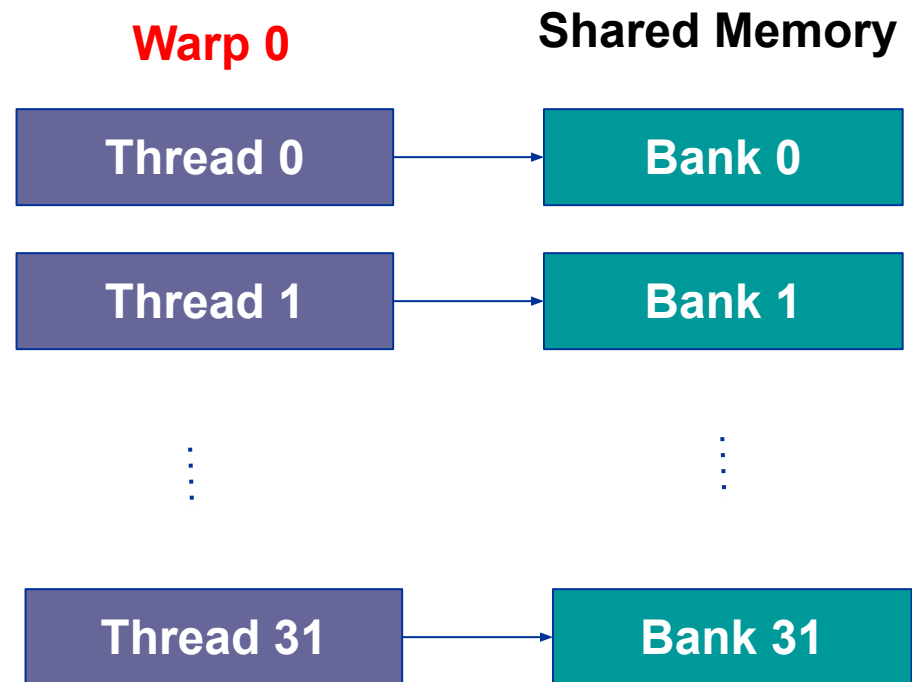


All are parallel!

Shared Memory Banks


- In the modern GPUs (> 2.x), successive 32-bit words map to successive banks.
 - ▣ $\text{bank} = (\text{address} / 4) \% 32$
- No. of banks = 32

```
__shared__ float shared[32];  
int S=1;  
float data = shared[S*tid];
```



All are parallel!

Shared Memory Banks




All are parallel!!

Shared Memory Bank Conflicts

- If shared memory address of any two threads falls in same bank (except when both of them access same address) then a *conflict* occurs
 - Access become *serial*


2-way bank conflict!



Shared Memory Bank Conflicts


- If shared memory address of any two threads falls in same bank (except when both of them access same address) then a *conflict* occurs
 - Access become *serial*

No conflict!




Exercise


- *Do the patterns have bank conflicts?*



(a)



(b)



(c)

Exercise-1:

- In the modern GPUs (> 2.x), successive 32-bit words map to successive banks.
 - ▣ $\text{bank} = (\text{address} / 4) \% 32$
- No. of banks = 32.
- *Does the following snippet have bank conflicts?*

```
__shared__ float shared[64];  
int S=2;  
float data = shared[S*tid];
```

Dynamic Shared Memory

- When the amount of shared memory required is unknown at compile-time, dynamic shared memory can be used.
- This is specified as the third parameter of kernel launch.

Dynamic Shared Memory

```
#include <stdio.h>
#include <cuda.h>

__global__ void dynshared() {
    extern __shared__ int s[];

    s[threadIdx.x] = threadIdx.x;
    __syncthreads();

    if (threadIdx.x % 2) printf("%d\n", s[threadIdx.x]);
}

int main() {
    int n;
    scanf("%d", &n);
    dynshared<<<1, n, n * sizeof(int)>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

Configurable L1 Cache and Shared Memory

- Shared memory and L1 cache can be configured by the programmer.
 - ❑ `cudaDeviceSetCacheConfig(kernelname, param);`
 - ❑ *kernelname* is the name of your kernel
 - ❑ *param*:
 - `cudaFuncCachePreferNone`: no preference for shared memory or L1 (default)
 - `cudaFuncCachePreferShared`: prefer larger shared memory and smaller L1 cache
 - `cudaFuncCachePreferL1`: prefer larger L1 cache and smaller shared memory
 - `cudaFuncCachePreferEqual`: prefer equal size L1 cache and shared memory

L1 Cache and Shared Memory

```
__global__ void dkernel() {  
    __shared__ unsigned data[BLOCKSIZE];  
    data[threadIdx.x] = threadIdx.x;  
}  
int main() {  
    cudaFuncSetCacheConfig(dkernel, cudaFuncCachePreferL1);  
    //cudaFuncSetCacheConfig(dkernel, cudaFuncCachePreferShared);  
    dkernel<<<1, BLOCKSIZE>>>();  
    cudaDeviceSynchronize();  
}
```

Thank you!